

Assignment 2: Multi-Task Hyperparameter Investigation

Team Members:

Abdulsalam Ashraf Aldwebi ID: 2210245306
Mohammed Ali Mayouf ID: 2190207878
Zidan Bagi

11.12.2025 (Fall 2025)

5.1 Front Matter

- **Total Training Time:** ≈ 21 hours
- **Total Successful Runs:** 15

5.2 Executive Summary

Best Hyperparameter Values

Hyperparameter	Best Value Found
Backbone	ResNet34
Segmentation Head	FCN
Activation	GELU
Initialization	Kaiming
Learning Rate	1e-2
Dropout Rate	0.2
Loss Function	Cross Entropy + Smooth L1

Table 1: Best value found for each hyperparameter across all experiments.

Overall Best Configuration

The best performing model was **Test 6 (act_gelu)** which achieved the highest mIoU and mAP.

- **Architecture:** ResNet34 + FCN + FPN
- **Activation:** GELU
- **Initialization:** Kaiming
- **Optimization:** LR=1e-3, Batch=16
- **Metrics:** mIoU: 13.7%, mAP: 15.6%, Pixel Acc: 75.5%

Top 3 Impactful Concepts

1. **Segmentation Head:** Switching from UNet to FCN provided the largest jump in performance ($8.9\% \rightarrow 13.5\%$ mIoU).
2. **Backbone Depth:** ResNet34 consistently outperformed ResNet18 ($13.5\% \text{ vs } 12.1\%$ mIoU).
3. **Activation Function:** GELU provided a noticeable improvement over ReLU and Leaky ReLU.

5.3 Core Analysis (Phase-by-Phase)

Phase 1: Network Architecture

1. Backbone & Head Selection

We compared ResNet18 vs ResNet34 and UNet vs FCN.




Figure 1: Comparison of Segmentation Heads. FCN (Right) shows lower validation loss and better convergence than UNet (Left).

Run ID	Backbone	Head	Train Loss	Val Loss	mIoU (%)	mAP (%)	Pixel Acc (%)
test_1	ResNet18	UNet	0.20	1.84	9.5%	12.1%	73.8%
test_2	ResNet34	UNet	0.34	1.84	8.9%	9.8%	72.5%
test_3	ResNet34	FCN	0.12	1.77	13.5%	14.8%	75.7%
test_4	ResNet18	FCN	0.11	1.86	12.1%	12.2%	75.1%

Table 2: Architecture Comparison Metrics (Final Epoch Values)

Discussion Convergence: ResNet34 with FCN converged faster and to a lower validation loss (1.77) than the UNet counterpart (1.84). The FCN head achieved the lowest training loss (0.12), indicating strong learning capacity. **Generalization:** The FCN model maintained better balance between train and validation performance, with the highest pixel accuracy (75.7%). **Recommendation:** We chose **ResNet34 + FCN** as the baseline for subsequent architectural experiments due to its superior mIoU (+4.6% over UNet) and pixel accuracy (+3.2%).

2. Activation Functions

We tested ReLU, Leaky ReLU, and GELU with the ResNet34+FCN architecture.




Figure 2: Activation Function Comparison.

Run ID	Activation	Train Loss	Val Loss	mIoU (%)	mAP (%)	Pixel Acc (%)
test_3	ReLU	0.12	1.77	13.5%	14.8%	75.7%
test_5	Leaky ReLU	0.74	1.64	11.2%	13.8%	68.0%
test_6	GELU	0.12	1.82	13.7%	15.6%	75.5%

Table 3: Activation Function Metrics (Final Epoch Values)

Discussion Convergence: GELU and ReLU both achieved low training loss (0.12), showing strong learning. Leaky ReLU struggled with higher training loss (0.74), indicating slower convergence. **Generalization:** Despite slightly higher validation loss, GELU achieved the best mAP (15.6%) and competitive pixel accuracy (75.5%). **Recommendation:** We selected **GELU** as it provided the highest mIoU and mAP, likely due to its smoother gradient flow helping the deep ResNet34 backbone.

3. Initialization Schemes

We tested Kaiming vs Xavier initialization with GELU activation.

Run ID	Init Scheme	Train Loss	Val Loss	mIoU (%)	mAP (%)	Pixel Acc (%)
test_6	Kaiming	0.12	1.82	13.7%	15.6%	75.5%
test_7	Xavier	0.10	1.81	12.6%	17.2%	74.8%
test_8	Normal	0.11	2.02	11.0%	13.3%	73.9%

Table 4: Initialization Scheme Metrics (Final Epoch Values)

Discussion Convergence: All schemes achieved similar training loss (0.10-0.12), but validation loss varied significantly. Normal initialization showed the highest validation loss (2.02). **Recommendation:** **Kaiming** initialization provided the best balance with highest mIoU, making it the preferred choice for ReLU-family activations.

Phase 2: Optimization (Learning Rate)

Using a ResNet34+UNet+GELU baseline (best_v1), we swept learning rates from 1e-5 to 1e-2.




Figure 3: Learning Rate Impact. Low LR (Left) learns too slowly. High LR (Right) converges fast but can be unstable.

Run ID	Learning Rate	Train Loss	Val Loss	mIoU (%)	mAP (%)	Pixel Acc (%)
lr_very_low	1e-5	0.70	2.25	7.3%	2.6%	70.6%
lr_low	1e-4	0.17	3.24	8.8%	7.6%	72.0%
best_v1	1e-3	0.42	1.57	9.1%	14.9%	73.4%
lr_high	1e-2	0.14	2.22	9.3%	14.3%	74.0%

Table 5: Learning Rate Metrics (Final Epoch Values)

Discussion Convergence: 1e-5 was too slow (high training loss 0.70 indicates underfitting). 1e-2 achieved the lowest training loss (0.14) but with higher validation loss (2.22), suggesting overfitting. 1e-3 provided the best balance. **Generalization:** The 1e-3 learning rate achieved the lowest validation loss (1.57) and best mAP (14.9%), indicating superior generalization. **Recommendation:** We recommend **1e-3** as the optimal learning rate, providing the best trade-off between convergence speed and generalization.

Phase 3: Regularization (Dropout)

We added Dropout to the baseline to combat overfitting.




Figure 4: Effect of Dropout on Loss Curves.

Run ID	Dropout Rate	Train Loss	Val Loss	mIoU (%)	mAP (%)	Pixel Acc (%)
best_v1	0.0	0.42	1.57	9.1%	14.9%	73.4%
drop_0.2	0.2	0.19	1.91	9.6%	12.4%	73.3%
drop_0.5	0.5	0.19	1.91	9.6%	12.4%	73.3%

Table 6: Regularization Metrics (Final Epoch Values)

Discussion Convergence: Dropout reduced training loss from 0.42 to 0.19, showing the model learned more robust features. However, validation loss increased slightly ($1.57 \rightarrow 1.91$). **Generalization:** Dropout improved segmentation mIoU from 9.1% to 9.6%, indicating better pixel-level generalization. Pixel accuracy remained stable at 73.3%. The slight drop in mAP suggests dropout may have regularized the detection head too aggressively. **Recommendation:** **Dropout 0.2** is recommended as a mild regularizer that improves segmentation without severely harming detection. Both 0.2 and 0.5 showed identical performance, suggesting 0.2 is sufficient.

Phase 4: Best Configuration Refinement

We combined the best hyperparameters discovered and tested two final configurations.

Run ID	Config	Train Loss	Val Loss	mIoU (%)	mAP (%)	Pixel Acc (%)
best_v1	R34+UNet+GELU	0.42	1.57	9.1%	14.9%	73.4%
best_v2	R34+FCN+GELU	1.48	1.57	4.6%	3.9%	73.2%

Table 7: Best Configuration Comparison (Final Epoch Values)

Discussion Interestingly, best_v2 showed poor performance despite combining optimal hyperparameters. The high training loss (1.48) suggests a training issue or incompatible hyperparameter combination. This highlights the importance of systematic experimentation rather than simply combining individually optimal values.

5.6 Visualization Appendix

[Insert your best/worst prediction images here manually in Overleaf]

5.7 Final Recommendations

Optimal Configuration

Based on our experiments, the optimal configuration is:

```
CONFIG = {
    "model": {
        "backbone": "resnet34",
        "segmentation_head": "fcn",
        "detection_head": "fpn",
        "activation": "gelu",
        "init_scheme": "kaiming",
        "dropout_rate": 0.2
    },
    "training": {
        "lr": 1e-3,
        "batch_size": 16,
        "seg_loss": "cross_entropy",
        "det_loss": "smooth_l1"
    }
}
```

Expected Performance:

- Training Loss: ~0.12
- Validation Loss: ~1.77
- mIoU: ~13.7%
- mAP: ~15.6%
- Pixel Accuracy: ~75.5%

Concept Ranking (1=Most Important)

1. **Segmentation Head Architecture** (UNet vs FCN) - Largest impact on mIoU (+4.6%)
2. **Backbone Capacity** (ResNet18 vs 34) - Significant capacity improvement
3. **Learning Rate** - Critical for convergence and generalization balance
4. **Activation Function** - GELU provided best mAP (+0.8% over ReLU)
5. **Initialization Scheme** - Kaiming best for ReLU-family activations
6. **Dropout Rate** - Mild regularization improved segmentation
7. Loss Function Choice - Standard choices worked well
8. Batch Size - 16 provided good stability
9. Weight Decay - Not extensively tested
10. Data Augmentation - Not tested in this study

Key Insights

- **Train-Val Gap:** Most models showed significant overfitting (train loss 0.1-0.4 vs val loss 1.5-2.2), suggesting stronger regularization or data augmentation could help.
- **Pixel Accuracy Plateau:** Pixel accuracy remained in 70-76% range across all experiments, suggesting architectural limitations or dataset difficulty.
- **Multi-Task Trade-off:** Improvements in segmentation (mIoU) sometimes came at the cost of detection (mAP), highlighting the challenge of multi-task optimization.

Future Work

1. **Transformer Backbone:** Try Swin Transformer or ViT for better global context and potentially higher pixel accuracy.
2. **DeepLabV3+ Head:** Implement DeepLabV3+ for segmentation to capture multi-scale context better than FCN.
3. **Cosine Annealing Scheduler:** Use a more advanced LR scheduler to converge to better minima and reduce train-val gap.
4. **Data Augmentation:** Test heavy augmentation to improve generalization and reduce overfitting.
5. **Multi-Task Loss Balancing:** Investigate dynamic loss weighting strategies to better balance segmentation and detection objectives.